

Get ready for Arm64 .. and an ARMful of container build cases

Syl Taylor

Specialist Solutions Architect

Compute

Ovidiu Valeanu

Specialist Solutions Architect

Containers

Arm-based compute options

On-premises and Edge

- Vendor servers (e.g. GIGABYTE)
- Apple MacBook with M1 chips
- Ampere Workstations
- Raspberry Pi ≥ 3
- Mobile phones & IoT devices
- Supercomputers (e.g. Fujitsu A64FX)



Cloud

- AWS Graviton, Ampere Altra, Alibaba Yitian, etc.
- Bare-metal, VMs, and managed services



Over 230 billion Arm chips produced

Running software on arm64

Applications

Libraries & Frameworks


Container Images

Container Runtime

Operating System

Hypervisor

**Entire software stack
must be compatible
with the CPU architecture (arm64)**



arm64 is also known as aarch64

Different CPU architectures

x86_64

```
mov    DWORD PTR [rbp-4], 10
mov    DWORD PTR [rbp-8], 20
mov    edx, DWORD PTR [rbp-4]
mov    eax, DWORD PTR [rbp-8]
add    eax, edx
```

arm64

```
mov    w0, 10
str    w0, [sp, 12]
mov    w0, 20
str    w0, [sp, 8]
ldr    w1, [sp, 12]
ldr    w0, [sp, 8]
add    w0, w1, w0
```

different ISA, different instructions

```
int sum() {
    int a = 10, b = 20;
    return a + b;
}
```

same code

-O0 (no optimizations)

Select a compatible base image

Use empty images: "FROM scratch"

OR

The screenshot shows the Docker Hub search results for the 'nginx' image. On the left, there are filters for 'Images', 'Extensions', and 'Plugins'. Under 'Trusted Content', 'Docker Official Image' is selected. Under 'Operating Systems', 'Linux' is selected. Under 'Architectures', 'ARM 64' is selected and highlighted with a red box. The search results on the right show three images: 'nginx' (Docker Official Image), 'bitnami/nginx', and 'bitnami/nginx-ingress-controller'. The 'nginx' image is highlighted with a red box around its architecture tags, which include 'ARM 64'.

☐ Images

☐ Extensions

☐ Plugins

Trusted Content

☐ Docker Official Image

☐ Verified Publisher

☐ Sponsored OSS

Operating Systems

☐ Linux

☐ Windows

Architectures

☐ ARM

☒ ARM 64

nginx DOCKER OFFICIAL IMAGE · 1B+ · 10K+

Updated 8 days ago

Official build of Nginx.

Linux **ARM 64** 386 mips64le PowerPC 64 LE IBM Z x86-64 ARM

bitnami/nginx VERIFIED PUBLISHER · 100M+ · 151

By VMware · Updated 5 days ago

Bitnami nginx Docker Image

Linux x86-64 arm64

bitnami/nginx-ingress-controller VERIFIED PUBLISHER · 10M+

By VMware · Updated 5 days ago

Bitnami Docker Image for NGINX Ingress Controller

Registries will use a tag
like **"ARM 64"**
with e.g. OS/ARCH:
"linux/arm64/v8"

Multi-arch container images

```
$ docker manifest inspect nginx:1.23.3
```

```
{
  "schemaVersion": 2,
  "mediaType": "application/vnd.docker.distribution.manifest.list.v2+json",
  "manifests": [
    {
      "mediaType": "application/vnd.docker.distribution.manifest.v2+json",
      "size": 1570,
      "digest": "sha256:942ae2dfd73088b54d7151a3c3fd5af038a51c50029bfcfd21f1e650d9579967",
      "platform": {
        "architecture": "amd64",
        "os": "linux"
      }
    },
    {
      "mediaType": "application/vnd.docker.distribution.manifest.v2+json",
      "size": 1570,
      "digest": "sha256:d415dd4e87d75f4d6607340d6d6fad78b0ed66b9c809eedb79ff08f3e58d008f",
      "platform": {
        "architecture": "arm64",
        "os": "linux",
        "variant": "v8"
      }
    }
  ]
}
```

Running an arm64 container image

On an **arm64** machine:

```
$ docker pull nginx:1.23.3
```

```
$ docker image inspect nginx:1.23.3 | grep 'Architecture'
"Architecture": "arm64",
```

```
$ docker run --name docker-nginx -p 127.0.0.1:1080:80 nginx:1.23.3
```

```
/docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt to perform configuration
```

```
...
```

```
/docker-entrypoint.sh: Configuration complete; ready for start up
```

```
...
```

```
2023/03/10 13:27:35 [notice] 1#1: nginx/1.23.3
```

```
2023/03/10 13:27:35 [notice] 1#1: built by gcc 10.2.1 20210110 (Debian 10.2.1-6)
```

```
2023/03/10 13:27:35 [notice] 1#1: OS: Linux 5.15.0-1031-aws
```

```
2023/03/10 13:27:35 [notice] 1#1: getrlimit(RLIMIT_NOFILE): 1048576:1048576
```

```
2023/03/10 13:27:35 [notice] 1#1: start worker processes
```

```
...
```

Watch out for: **exec format error**

On an **arm64** machine:

```
$ docker run --platform linux/amd64 nginx:1.23.3
```

Unable to find image 'nginx:1.23.3' locally

1.23.3: Pulling from library/nginx

3f9582a2cbe7: Already exists

9a8c6f286718: Pull complete

e81b85700bc2: Pull complete

73ae4d451120: Pull complete

6058e3569a68: Pull complete

3a1b8f201356: Pull complete

Digest: sha256:aa0afebbb3cfa473099a62c4b32e9b3fb73ed23f2a75a65ce1d4b4f55a5c2ef2

Status: Downloaded newer image for nginx:1.23.3

exec /docker-entrypoint.sh: exec format error

Building arm64 container images

Method #1: Natively (on an **arm64** machine, *locally or remote builds*)

```
$ docker build -t example:0.1 .
```

Method #2: Emulation (on an **x86_64** machine)

```
$ docker buildx build --platform linux/arm64 .
```

Method #3: Cross-compilation (on an **x86_64** machine)

```
FROM --platform=linux/amd64 ubuntu as build  
# RUN <install build dependencies>  
# COPY <source> .  
# RUN <cross-compiler for target: linux/arm64, -o /out/arm64_binary>
```

```
FROM --platform=linux/arm64 ubuntu as runtime  
# RUN <install runtime dependencies via emulation>  
COPY --from=build /out/arm64_binary /bin
```

Buildah and multi-arch images

- **Daemon-less, rootless, OCI multi-arch compatible image builder**
- **Build container from scratch or pre-existing Dockerfile**

```
# Create a multi-architecture manifest
buildah manifest create ${MANIFEST_NAME}
```

```
# Build your amd64 architecture container
buildah bud \ --tag "${REGISTRY}/${USER}/${IMAGE_NAME}:${IMAGE_TAG}" \ --
manifest ${MANIFEST_NAME} \ --arch amd64 \ ${BUILD_PATH}
```

```
# Build your arm64 architecture container
buildah bud \ --tag "${REGISTRY}/${USER}/${IMAGE_NAME}:${IMAGE_TAG}" \ --
manifest ${MANIFEST_NAME} \ --arch arm64 \ ${BUILD_PATH}
```

```
# Push the full manifest, with both CPU architectures
```

```
buildah manifest push --all \ ${MANIFEST_NAME} \
"docker://${REGISTRY}/${USER}/${IMAGE_NAME}:${IMAGE_TAG}"
```

Emulation sounds great! But is it?

Let's build an **arm64** container image from machines with 16 vCPUs & 32 GiB

```
$ git clone https://github.com/IntelRealSense/librealsense.git
$ cd librealsense/scripts/Docker
$ LIBRS_GIT_TAG=`git describe --abbrev=0 --tags`
$ LIBRS_VERSION=${LIBRS_GIT_TAG#"v"}
```

Emulated build – from **x86_64**

```
$ time docker buildx build --platform linux/arm64 -t librealsense-arm64:0.1 . --build-arg LIBRS_VERSION=$LIBRS_VERSION
... 22m27.874s...
```

Native build – on **arm64**

```
$ time docker build -t librealsense-arm64:0.1 . --build-arg LIBRS_VERSION=$LIBRS_VERSION
... 3m1.891s
```

Emulation is slow

Let's look at **Dockerfile**

```
....  
RUN apt-get update && \  
    apt-get install -qq -y --no-install-recommends build-essential cmake git libssl-dev libusb- 1.0-0-dev pkg-config libgtk-3-dev  
libglfw3-dev libgl1-mesa-dev libglu1-mesa-dev curl python3 python3-dev ca-certificates && \  
....  
RUN cd /usr/src/librealsense && \  
    mkdir build && cd build && \  
    cmake \  
    -DCMAKE_C_FLAGS_RELEASE="${CMAKE_C_FLAGS_RELEASE} -s" \  
    -DCMAKE_CXX_FLAGS_RELEASE="${CMAKE_CXX_FLAGS_RELEASE} -s" \  
    -DCMAKE_INSTALL_PREFIX=/opt/librealsense \  
    -DBUILD_GRAPHICAL_EXAMPLES=OFF \  
    -DBUILD_PYTHON_BINDINGS:bool=true \  
    -DCMAKE_BUILD_TYPE=Release ../ && \  
    make -j$(($(nproc)-1)) all && \  
    make install  
...
```

Note the **RUN** commands which during emulation will require instruction translation

The image features a dark gray background with a subtle pattern of concentric circles. In each of the four corners, there are decorative elements resembling circuit board traces or neural network connections, consisting of thin white lines and small white circles.

But if we stick to native builds, then..

Case 1: No changes needed (**build**)

Dockerfile

```
FROM public.ecr.aws/lts/ubuntu:20.04_stable

WORKDIR /home/app
COPY requirements.txt .

RUN apt update && \
    apt install -y --no-install-recommends python3 python3-pip && \
    pip install -r requirements.txt
```

requirements.txt

```
numpy==1.21.6
scikit-learn==1.2.0
matplotlib==3.6.1
```

```
$ docker build -t sklearn-example:0.1 .
```

Case 1: No changes needed (runtime)

app.py

```
import numpy as np
from sklearn.tree import DecisionTreeRegressor
import matplotlib.pyplot as plt


# Create a random dataset
rng = np.random.RandomState(1)
...

# Fit regression model
...
regr_1.fit(X, y)
regr_2.fit(X, y)

# Predict
...
y_1 = regr_1.predict(X_test)
y_2 = regr_2.predict(X_test)

# Plot the results
...
plt.savefig("plot.png")
```

plot.png



```
$ docker run -v ./:/home/app -it sklearn-example:0.1 /usr/bin/python3 app.py
```



Case 1: No changes needed (**demo**)

View file: [demo-case-1.mp4](#)



Case 1: Why are no changes needed?

Move from x86_64/amd64 -> arm64 with **no effort**:

1. `public.ecr.aws/lts/ubuntu:20.04_stable` is a **multi-arch base image**
2. Python *pip wheels* & versions are available for both **x86_64** & **arm64**
3. Packages installed with *apt* are available for both **x86_64** & **arm64**
4. `app.py` is written in **pure Python** (no hardware-specific low-level code)

How about the other cases?

Watch out for:

- **Successful builds are great, but also check for runtime errors**

Build success. Runtime error.

```
> [3/3] RUN python3 -c 'import tvn; print(tvn.__version__)':  
...  
#6 0.587 ModuleNotFoundError: No module named 'typing_extensions'  
-----  
Dockerfile:7  
-----  
5 |         pip install apache-tvm  
6 |  
7 | >>> RUN python3 -c 'import tvn; print(tvn.__version__)'  
8 |  
-----
```

- **Validate your applications through software testing**
(for example, CPU architecture changes can affect floating number comparisons)

Case 2: Old version issue

Dockerfile

```
FROM public.ecr.aws/lts/ubuntu:20.04_stable

WORKDIR /home/app

RUN apt update && \
    apt install -y --no-install-recommends python3-pip

RUN pip install numpy==1.18.5
```

app.py

```
import numpy as np
print(np.__version__)
print(np.arange(15).reshape(3, 5))
```

Case 2: Old version issue

Error – arm64

```
#0 11.84      raise RuntimeError("Broken toolchain: cannot link a simple C program")
#0 11.84      RuntimeError: Broken toolchain: cannot link a simple C program
```

x86_64

```
#7 [4/4] RUN pip install numpy==1.18.5
#7 1.548 Collecting numpy==1.18.5
#7 1.563   Downloading numpy-1.18.5-cp38-cp38-
manylinux1_x86_64.whl (20.6 MB)
#7 2.141 Installing collected packages: numpy
#7 3.538 Successfully installed numpy-1.18.5
#7 DONE 3.8s
```

arm64

```
#7 [4/4] RUN pip install numpy==1.18.5
#7 1.278 Collecting numpy==1.18.5
#7 1.295   Downloading numpy-1.18.5.zip (5.4 MB)
#7 1.587   Installing build dependencies: started
#7 4.239   Installing build dependencies: finished with
status 'done'
#7 4.242   Getting requirements to build wheel: started
#7 4.448   Getting requirements to build wheel: finished
with status 'done'
#7 4.450   Preparing wheel metadata: started
#7 11.54   Preparing wheel metadata: finished with
status 'error'
```

no pre-built pip wheel available

Case 2: Old version issue (*ideas*)

Potential solutions to fix the pip wheel for numpy on arm64:

1. Install *build dependencies* and version *1.18.5* will work

```
RUN apt update && \  
    apt install -y --no-install-recommends gpg-agent software-properties-common python3-pip build-essential && \  
    add-apt-repository -y ppa:deadsnakes/ppa && \  
    apt install -y python3.7 python3.7-dev python3.7-distutils && \  
    update-alternatives --install /usr/bin/python3 python3 /usr/bin/python3.7 1
```

2. Keep the current Dockerfile, but install *version >= 1.19.0*
3. Use different package repo like *apt install python3-numpy (1.17.4)*

Case 2: Old version issue (demo)

View file: [demo-case-2.mp4](#)

Case 3: No binary available

Dockerfile

```
FROM public.ecr.aws/lts/ubuntu:20.04_stable

RUN apt update && \
    apt install -y --no-install-recommends python3 python3-pip && \
    pip install apache-tvm==0.9.0 typing-extensions

RUN python3 -c 'import tvm; print(tvm.__version__)'
```

Error – arm64

```
#5 9.707 ERROR: Could not find a version that satisfies the requirement apache-tvm==0.9.0 (from versions: none)
#5 9.707 ERROR: No matching distribution found for apache-tvm==0.9.0
```

Case 3: No binary available (build)

Change Dockerfile. **Build from source**. An example below:

```
FROM public.ecr.aws/lts/ubuntu:20.04_stable

ARG LLVM_VERSION="llvmorg-14.0.6"
ARG LLVM_BIN="clang+llvm-14.0.6-aarch64-linux-gnu"
ARG TVM_VERSION="v0.9.0"
ARG DLPACK_VERSION="v0.5"

RUN apt update && \
    apt install -y --no-install-recommends git wget python3 python3-pip python3-dev python3-setuptools && \
    apt install -y --no-install-recommends libtinfo-dev zlib1g-dev build-essential cmake libedit-dev libxml2-dev libncurses5 && \
    cd /tmp; wget https://github.com/llvm/llvm-project/releases/download/${LLVM_VERSION}/${LLVM_BIN}.tar.xz && \
    mkdir llvm; tar -xf ${LLVM_BIN}.tar.xz -C /tmp/llvm && \
    git clone --recursive https://github.com/apache/tvm tvm && \
    cd tvm; git checkout tags/${TVM_VERSION} && \
    cd 3rdparty/dlpack; git checkout tags/${DLPACK_VERSION}; cd ../.. && \
    mkdir build && \
    sed -i "/set(USE_LLVM OFF)/c\\set(USE_LLVM /tmp/llvm/${LLVM_BIN}/bin/llvm-config)" cmake/config.cmake && \
    cp cmake/config.cmake build && \
    cd build; cmake ..; make -j$(nproc); cd .. && \
    cd python; python3 setup.py install --user; cd ..

RUN python3 -c 'import tvm; print(tvm.__version__)'
```


Case 3: No binary available (demo)

View file: [demo-case-3.mp4](#)

Case 4: Dependency issue

Dockerfile

```
FROM public.ecr.aws/lts/ubuntu:20.04_stable

RUN apt update && \
    apt install -y --no-install-recommends python3 python3-pip && \
    pip install confluent_kafka

RUN python3 -c 'import confluent_kafka; print(confluent_kafka.version())'
```

Error 1 – arm64 (trying to build from source)

```
#5 9.868 Downloading confluent-kafka-2.0.2.tar.gz (119 kB)
...
unable to execute 'aarch64-linux-gnu-gcc': No such file or directory
```

Error 2 – arm64 (dependency with specific version not found)

```
#5 21.86 66 | #error "confluent-kafka-python requires librdkafka v2.0.2 or later.
```

Case 4: Dependency issue (build)

Change Dockerfile. **Build** both the **pip wheel** and **librdkafka from source**. An example:

```
FROM public.ecr.aws/lts/ubuntu:20.04_stable

ARG LIBRDKAFKA_VERSION="2.0.2"

RUN apt update && \
    apt install -y --no-install-recommends git build-essential python3 python3-pip python3-dev && \
    git clone -b v${LIBRDKAFKA_VERSION} https://github.com/confluentinc/librdkafka && \
    cd librdkafka && ./configure --install-deps && \
    make && \
    make install && \
    ldconfig && \
    pip install confluent-kafka==2.0.2 --no-binary :all:

RUN python3 -c 'import confluent_kafka; print(confluent_kafka.version())'
```

Case 4: Dependency issue (demo)

View file: [demo-case-4.mp4](#)

Case 5: Application issue

Dockerfile

```
FROM public.ecr.aws/lts/ubuntu:20.04_stable

WORKDIR /home/app
# Install certificates for github and proxy.golang

RUN apt update && \
    apt install -y --no-install-recommends wget git && \
    wget --no-check-certificate https://go.dev/dl/go1.20.2.linux-amd64.tar.gz && \
    rm -rf /usr/local/go && tar -C /usr/local -xzf go1.20.2.linux-amd64.tar.gz && \
    export PATH=$PATH:/usr/local/go/bin && \
    git clone https://github.com/segmentio/parquet-go.git && \
    cd parquet-go && \
    go test -tags amd64 -v ./
```

Error – arm64

```
...go/src/crypto/internal/boring/sig/sig_amd64.s:38: unrecognized instruction "BYTE"
...go/.../klauspost/compress@v1.15.9/internal/cpuinfo/cpuinfo_amd64.s:25: unrecognized instruction "BTQ"
...go/.../klauspost/compress@v1.15.9/internal/cpuinfo/cpuinfo_amd64.s:26: unrecognized instruction "SETCS"
#10 29.64 asm: too many errors ...
```

Case 5: Application issue

Thanks to tags (e.g. `//go:build purego || !amd64`), we can still use package without the asm.

Change Dockerfile. **Test purego** instead of **amd64** (which has non-arm64 assembly .s files) .

```
FROM public.ecr.aws/lts/ubuntu:20.04_stable

WORKDIR /home/app

# Install certificates for github and proxy.golang

RUN apt update && \
    apt install -y --no-install-recommends wget git && \
    wget --no-check-certificate https://go.dev/dl/go1.20.2.linux-amd64.tar.gz && \
    rm -rf /usr/local/go && tar -C /usr/local -xzf go1.20.2.linux-amd64.tar.gz && \
    export PATH=$PATH:/usr/local/go/bin && \
    git clone https://github.com/segmentio/parquet-go.git && \
    cd parquet-go && \
    go test -tags purego -v ./
```

Case 5: Application issue (demo)

View file: [demo-case-5.mp4](#)

Other notable cases

- **TensorFlow** for arm64 is not built with CUDA support for Nvidia GPUs
- **PostgreSQL** should be compiled with Arm's LSE (atomic instructions)
- **Emscripten** build failed when a new package was added (cipd – no linux-arm64)
- **Go compiler** binary, e.g. goX.linux-**amd64**.tar.gz -> goX.linux-**arm64**.tar.gz
- Re-compiling with **Go** \geq **1.18** leads to up to 20% improvement on arm64

Ecosystem

The state of the ecosystem is **continuously changing**, meaning more **arm64 binaries** and **optimizations** are added as we speak...


As a general rule: newer software versions tend to have arm64 support


Open-source call to action

Tell software maintainers when you need **arm64/aarch64** binaries or container images.

Open a new **Issue** or **+1** to an existing one asking for support like in this example:

Multi-arch deployment of Harbor is not supported (exec format error) #18242

 Open samip5 opened this issue last month · 3 comments




samip5 commented last month · edited

Expected behavior and actual behavior:
I tried to install Harbor to an arm64 instance, and it fails because the prepare container is only amd64. I would have expected it to not be an issue.

Steps to reproduce the problem:

1. Create a Oracle Cloud ampere instance or AWS Gaviton
2. Try to deploy Harbor
3. Have it fail on install due to:

```
[Step 3]: preparing harbor configs ...
prepare base dir is set to /home/ubuntu/harbor
WARNING: The requested image's platform (linux/amd64) does not match the
exec /usr/bin/python3: exec format error
```

Assignees
 wy65701436

Labels
None yet

Projects
None yet

Milestone
No milestone

Development
No branches or pull requests